

# Recap: Robust MPC

- Robust MPC solves the following Optimal Control problem as follows:

$$\begin{aligned}
 J_0(\bar{x}_0) = \min_{X, U, Z} \quad & \hat{J}_N(\bar{x}_N) + \sum_{i=0}^{N-1} \bar{x}_i^\top Q \bar{x}_i + u_i^\top R u_i \\
 \text{s.t.} \quad & \bar{x}_{i+1} = A\bar{x}_i + Bu_i, \quad \forall i \in \{0, 1, 2, \dots\} \\
 & \bar{x}_i \in \mathcal{X} \ominus \mathcal{R}_i, \quad \forall i \in \{1, 2, \dots\} \\
 & u_i = K\bar{x}_i + z_i, \quad \forall i \in \{0, 1, 2, \dots\} \\
 & u_i \in \mathcal{U} \ominus (K \circ \mathcal{R}_k), \quad \forall i \in \{0, 1, 2, \dots\} \\
 & x_0 = \bar{x}_0 \\
 & x_N \in \mathcal{X}_f \ominus \mathcal{R}_N
 \end{aligned}$$

The diagram includes several annotations with red arrows pointing to specific parts of the equations:
 

- Nominal stage cost** points to the term  $\bar{x}_i^\top Q \bar{x}_i$  in the cost function.
- Nominal dynamics** points to the equation  $\bar{x}_{i+1} = A\bar{x}_i + Bu_i$ .
- $z_i$  is effectively the control decision** points to the equation  $u_i = K\bar{x}_i + z_i$ .
- Robust Invariant Set** points to the terminal constraint  $x_N \in \mathcal{X}_f \ominus \mathcal{R}_N$ .

A green bracket on the left side of the constraints is labeled **Constraint Robustification**.

# Adding Learning to Robust MPC

- We will now add a “learning-based” component to our model:
  - (1) The goal is to “reduce” the conservativeness of our robust formulation
  - (2) We can use Machine Learning to improve our dynamics model
- Let’s restart by recalling the linear dynamics:

$$x_{k+1} = Ax_k + Bu_k + w_k$$

- We will add some non-linear function approximation  $h(x_k, u_k)$  in order to approximate  $w_k$ 
  - That way we can try to “predict” the effect of uncertainty before it is manifested
- Then the dynamics would then become:

$$x_{k+1} = Ax_k + Bu_k + h(x_k, u_k)$$

# Learning Dynamics

- Another interpretation for  $h(x_k, u_k)$  is to capture the “complex” components of the system dynamics:

$$\underbrace{x_{k+1} - Ax_k + Bu_k}_{\text{(linear) model mismatch}} = w_k$$

- So  $h(x_k, u_k)$  is an approximation of this mismatch.
- In this case, we say, the linear model is the **nominal model**:

$$\bar{x}_{k+1} = A\bar{x}_k + Bu_k$$

- In many applications the function  $h(x_k, u_k)$  is also called the **Oracle Function**.

# Learning-Based Model Predictive Control (LBMPC)

- As always, the LBMPC solves a N-step lookahead in a receding horizon fashion:

$$\begin{aligned} J_0(\bar{x}_0) = \min_{X, U, Z} \quad & \hat{J}_N(\tilde{x}_N) + \sum_{i=0}^{N-1} \tilde{x}_i^\top Q \tilde{x}_i + u_i^\top R u_i \\ \text{s.t.} \quad & \tilde{x}_{i+1} = A\tilde{x}_i + Bu_i + h(\tilde{x}_i, u_i), \quad \forall i \in \{0, 1, 2, \dots\} \\ & \bar{x}_{i+1} = A\bar{x}_i + Bu_i, \quad \forall i \in \{0, 1, 2, \dots\} \\ & u_i = K\bar{x}_i + z_i, \quad \forall i \in \{0, 1, 2, \dots\} \\ & \bar{x}_i \in \mathcal{X} \ominus \mathcal{R}_i, u_i \in \mathcal{U} \ominus (K \circ \mathcal{R}_k), \quad \forall i \in \{0, 1, 2, \dots\} \\ & x_0 = \bar{x}_0 \\ & \bar{x}_N \in \mathcal{X}_f \ominus \mathcal{R}_N \end{aligned}$$

# Learning-Based Model Predictive Control (LBMPC)

- Note that the key distinction here is that we keep two models:

**Nominal Model**

$$\bar{x}_{k+1} = A\bar{x}_k + Bu_k$$

- Enforce robust constraints.
- So, for any possible mismatch, the true system remains feasible:

$$x_{k+1} = A\bar{x}_k + Bu_k + w_k \in \mathcal{X}$$

**Learned Model**

$$\tilde{x}_{k+1} = A\tilde{x}_k + Bu_k + h(\tilde{x}_k, u_k)$$

- No constraints enforced on the learned model.
- The objective function is based on the learned model:

$$\hat{J}_N(\tilde{x}_N) + \sum_{i=0}^{N-1} \tilde{x}_i^\top Q \tilde{x}_i + u_i^\top R u_i$$

# LBMPC Properties

- It turns out that we can still ensure the two main properties, even with the Oracle:
  - (1) Recursive Feasibility
  - (2) (Robust) Asymptotic Stability
- We need that  $\mathcal{X}_f$  to be a robust control invariant set, associated with the LQR control law.
- Our terminal cost function approximation will be given as before:

$$\hat{J}_N(\bar{x}_N) = \bar{x}_N^\top P \bar{x}_N$$

- Lastly the function  $h_k(x_k, u_k)$  needs to be continuous and we need to satisfy:

$$h(x_k, u_k) \in \mathcal{W} \quad \longrightarrow \quad \left\{ \begin{array}{l} \text{So the approximation needs to be bounded and} \\ \text{be contained into the uncertainty polytope} \end{array} \right.$$

# LBMPC Properties

- The proofs follow the lines we covered last lecture:
  - Check the paper: “Provably Safe and Robust Learning-Based Model Predictive Control, Aswani et al. 2012” for the full proofs.
- Instead, we will focus our analyzes to designing an appropriate oracle function.
- For the MPC results to hold the function  $h(x_k, u_k)$  needs to be:
  - (1) Continuous
  - (2) Bounded
- We will add (3) differentiable to the requirements (although it is not needed for the proofs)
- The reason to add differentiable is to allow numerical solvers to take gradients in order to solve the optimal control problem.

# Parametric Oracles

- The first type of Oracles are the typical parametric functions:
  - The function  $h(x_k, u_k)$  is defined by a parameter vector  $\theta_k$
- Suppose we were able to run the system in a simulator, applying some control sequence and obtaining the following trajectory:

$$(x_0, u_0, x_1, u_1, \dots, x_{N-1}, u_{N-1}, x_N)$$

- Then we obtaining our oracle function, by solving the typical Regression Problem:

$$\theta^* \in \arg \min_{\theta} \left\{ \sum_{i=0}^{N-1} (Y_i - h(x_i, u_i; \theta))^2 \right\}$$

- Where:

$$Y_i = x_{i+1} - (Ax_i + Bu_i)$$



# Example: Linear Oracles

- The oracle function can be itself a linear function:

$$h(x_k, u_k; \theta_k) = F_k x_k + G_k u_k$$

- In this case, the oracle has a nice interpretation as being a correction to the dynamics


- The learned model then becomes:

$$\tilde{x}_{k+1} = (A + F_k)\tilde{x}_k + (B + G_k)u_k$$

- Where  $\theta_k = (F_k, G_k)$ . So we can see that oracle corrects the linear nominal model by essentially “updating” the matrices A and B.

# Example: DNN's

- The oracle function can be given by a Deep Neural Network (DNN):



- And the DNN would be training in a similar fashion as we saw for model-free methods.
- Hence, LBMPC algorithm would alternate between two steps:
  - (1)** Prediction step: updating the oracle
  - (2)** Feedback step: Solving the optimal control problem

# Non-Parametric Oracles

- The second type of Oracles are the ones that do not rely on parameters.
- One of such oracles is a kernel-based oracle called the **Nadararya-Watson Oracle**:

$$h(x_k, u_k) = \frac{\sum_{i=0}^{N-1} (x_{i+1} - Ax_i - Bu_i) \mathcal{K}\left(h^{-2} \left\| \begin{bmatrix} x_k \\ u_k \end{bmatrix} - \begin{bmatrix} x_i \\ u_i \end{bmatrix} \right\|^2\right)}{\lambda + \sum_{i=0}^{N-1} \mathcal{K}\left(h^{-2} \left\| \begin{bmatrix} x_k \\ u_k \end{bmatrix} - \begin{bmatrix} x_i \\ u_i \end{bmatrix} \right\|^2\right)}$$

- Where  $\mathcal{K}(\cdot)$  is a kernel function, and  $\lambda$  is some regularization hyper-parameter

# Example: Kernel-based Oracles

- The idea of using kernels in LBMPC is very nice, because it does not assume any prior form for the un-modelled dynamics.
- So the oracle essentially computes an weighted average based on the kernels:

$$h(x_k, u_k) = \frac{\sum_{i=0}^{N-1} Y_i \mathcal{K}(z_i)}{\lambda + \sum_{i=0}^{N-1} \mathcal{K}(z_i)}$$

- Where


$$Y_i = x_{i+1} - (Ax_i + Bu_i) \quad \mathcal{K}(z_i) = \mathcal{K}\left(h^{-2} \left\| \begin{bmatrix} x_k \\ u_k \end{bmatrix} - \begin{bmatrix} x_i \\ u_i \end{bmatrix} \right\|^2\right)$$

# Example: Kernel-based Oracles

- For example we can use Gaussian Kernels:

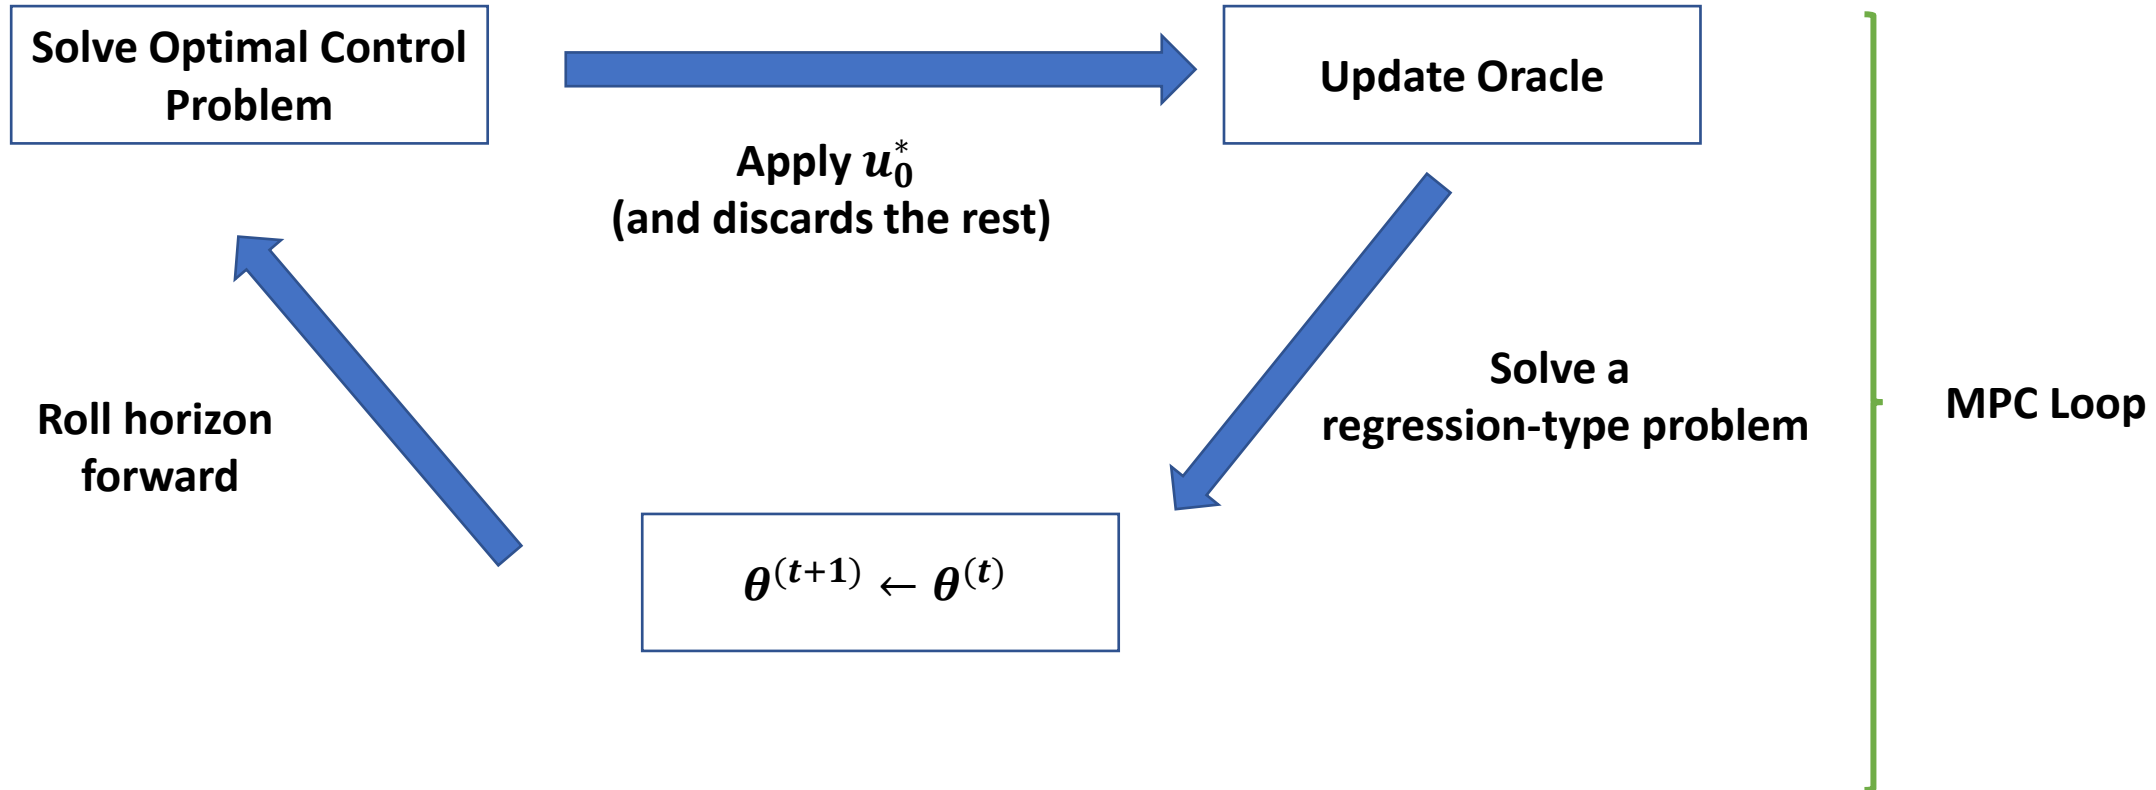
$$\mathcal{K}(z_i) = \frac{1}{\sqrt{2\pi}} \exp \left( - \left\| \begin{bmatrix} x_k \\ u_k \end{bmatrix} - \begin{bmatrix} x_i \\ u_i \end{bmatrix} \right\|^2 \right)$$

- We can represent this weighted-average by a figure:




# Learning-Based Model Predictive Control (LBMPC)

- We can represent the LBMPC in the following scheme:



# Example of LBMPC: Quadrotor Flight Control

- We illustrate an application of LBMPC with a Ball-Catching experiment by a quadrotor:



# Quadrotor Model

- The quadrotor drone can be modelled a linear system where the states are 3D-position, their time-derivatives, the rotation angles and their derivatives.
  - $(x_N, x_E, x_D)$  are the positions
  - $(\psi, \theta, \phi)$  are the rotation angles (yaw-pitch-roll).
  - The rotation  $\psi$  angle is held fix, for reference.
- It is common to work with two sets of reference frames:
  - (1) body-fixed frame
  - (2) inertial frame
- For ease of presentation, let's abstract the reference frames and just focus on the resulting linear system.
  - For a full description of the underlying physics we refer to: "Learning-Based Model Predictive Control on a Quadrotor: Onboard Implementation and Experimental Results. Bouffard et al."



# Quadrotor Model

- The linear dynamics are obtained by discretization of the continuous system, with steps  $\Delta t = 0.025s$ . And are given as follows, each horizontal axis:

$$x_{k+1} = Ax_k + Bu_k = \begin{bmatrix} 1 & 0.025 & 0.003 & 0 \\ 0 & 1 & 0.245 & 0 \\ 0 & 0 & 0.797 & 0.023 \\ 0 & 0 & -1.798 & 0.977 \end{bmatrix} x_k + \begin{bmatrix} 0 \\ 0 \\ 0.01 \\ 0.9921 \end{bmatrix} u_k$$

- For the vertical axis there is the effect the effect of gravity which changes the system matrices. For that component the linear dynamic is given by:

$$x_{k+1} = Ax_k + Bu_k = \begin{bmatrix} 1 & 0.025 \\ 0 & 1 \end{bmatrix} x_k + c_T \begin{bmatrix} 0.0003 \\ 0.025 \end{bmatrix} u_k + b_z$$

# Quadrotor Model

- The full system is defined by the concatenation of the three axis:

$$A = \text{blkdiag}(A, A, A_z) \in \mathbb{R}^{10 \times 10}$$

$$B = \text{blkdiag}(B, B, B_z) \in \mathbb{R}^{10 \times 3}$$

$$b = \begin{bmatrix} 0 & 0 & b_z \end{bmatrix}^\top$$

- And the **nominal system** evolves as:

$$\bar{x}_{k+1} = A\bar{x}_k + Bu_k + b$$

- The control inputs  $u_k$  are the thrust executed along each axis.

# Quadrotor Model

- Since we are using a linear model to approximate the flight dynamics we will add a linear oracle. The oracle will be time-varying in order to correct for mismatches in the linear model:

$$h(x_k, u_k) = Fx_k + Gu_k + v$$

- We can interpret this as  $(A, B)$  being the linearization (e.g.: derivate) information along some reference pair  $(x, u)$ . And  $(F, G)$  are correction when we move from  $(x, u)$  to  $(x_k, u_k)$ .
- Then the **learned model** becomes:

$$\tilde{x}_{k+1} = (A + F)\tilde{x}_k + (B + G)u_k + b + v$$

# Ball in free flight model

- The ball model is taken to represent the dynamics of ball being thrown in the air by a human.
- The ball falls due to gravity and spins, due to the human throwing it.
- We consider a linear model as well that incorporates air drag suffered by the ball while it is flight:

$$x_{k+1} = \text{blkdiag}(A_b, A_b, A_b)x_k + b_F$$

- Where  $b_F$  is an empirical offset vector that is dependent on air resistance

- And  $A_b$  is a double integrator, discretized for each axis:  $A_b = \begin{bmatrix} 1 & 0.025 \\ 0 & 1 \end{bmatrix}$

# Training the Oracle

- In the practical experiments, the authors used onboard sensors to estimate the quadrotor and the ball positions.
- That means that the state  $x_k$  is not fully observed.
- In particular, the sensors are only able to estimate the actual positions and angles, with the associate derivatives not being observable.
- So we can write the sensor information as:

$$y_k = Cx_k + \epsilon$$

- Where  $\epsilon$  is some white noise vector.

# Training the Oracle

- In a practical situation such as this. We need to resort to system identification techniques in order to infer the state values from the observations.
- One such technique is the Extended Kalman Filter (EKF).
  - The reference paper implements this in their practical experiments.
- We have not covered this topic in the course. So for this presentation let's suppose we do have the full system state observation.
  - That is, we are able to fully compute the nominal state  $\bar{x}_k$ .
- Then in this case, we can estimate the oracle parameters via the typical regression step:

# Training the Oracle

- The regression step is then as follows:

$$\theta^* \in \arg \min_{\theta} \left\{ \sum_{i=0}^{N-1} (Y_i - h(x_i, u_i; \theta))^2 \right\}$$

- Where  $\theta = (F, G, v)$  and:

$$Y_i = x_{i+1} - (Ax_i + Bu_i)$$

- Note that as the system progress we keep adding more and more “data points” to this regression problem.
  - We can let the quadrotor fly in many simulation runs in order to collect a sizeable data set of transitions  $(x_i, u_i, x_{i+1})$ .

# LBMPC for the quadrotor

- The LBMPC problem for the quadrotor is as follows:

$$\begin{aligned} J_0(\bar{x}_0, \theta_0) = \min_{X, U, Z} \quad & (\tilde{x}_N - x_s)^\top P(\tilde{x}_N - x_s) + \sum_{i=0}^{N-1} (\tilde{x}_i - x_s)^\top Q(\tilde{x}_i - x_s) + (u_i - u_s)^\top R(u_i - u_s) \\ \text{s.t.} \quad & \tilde{x}_{i+1} = (A + F_0)\tilde{x}_i + (B + G_0)u_i + b + v_0, \quad \forall i \in \{0, 1, 2, \dots\} \\ & \bar{x}_{i+1} = A\bar{x}_i + Bu_i, \quad \forall i \in \{0, 1, 2, \dots\} \\ & u_i = K\bar{x}_i + z_i, \quad \forall i \in \{0, 1, 2, \dots\} \\ & \bar{x}_i \in \mathcal{X} \ominus \mathcal{R}_i, u_i \in \mathcal{U} \ominus (K \circ \mathcal{R}_k), \quad \forall i \in \{0, 1, 2, \dots\} \\ & x_0 = \bar{x}_0 \\ & \bar{x}_N \in \mathcal{X}_f \ominus \mathcal{R}_N \end{aligned}$$

- Where as usual,  $\mathcal{X}_f$  is a robust control-invariant set associated with the LQR version of the problem. And the constraints are polyhedral sets



# LBMPC for the quadrotor

- The state-control pair  $(x_s, u_s)$  used as a reference trajectory are the desired set-point of the quadrotor:
  - $x_s$  is the predicted landing location for the ball.
  - $u_s$  is the control that keeps the quadrotor stationary at  $x_s$

- We can compute  $u_s$  by solving the following system of equations:

$$x_s = (A + F)x_s + (B + G)u_s + b + v_k$$

- Note that  $u_s$  may not be a feasible control. It does not need to be, it is only taken as a reference.
- We use the learned model in order to obtain the set-point reference.

# Experiments with LBMPC


- The authors implemented the LBMPC algorithm in an onboard computer:
  - 1.6GHz Intel Atom N260 CPU
  - 1 GB of Ram
  - WiFi communications to Vicon MX motion capture system to estimate vehicle and ball positions
- The planning horizon for the MPC is  $N = 15$ .
- Commands are issued at the rate of 40Hz.
- The optimal control problem faced by the LBMPC at every planning stage is a Quadratic Program (so quadratic objective, linear constraints).

# Experiments with LBMPC


- There were two main sets of experiments:
  - (1) Flight close to the ground
  - (2) Flight to an alternating set-point reference
- The first experiment is interesting because it is designed to show how the oracle can “learn” the aerodynamical effect that ground has on the quadrotor:
  - If the vehicle hover very close to the ground, it subject to additional lift due to the air being “reflected” back to the vehicle.
- It is very important effect when considering “soft landing”:
  - Landing smoothly without turning off the engines.

# Ground effects experiment

- The results can be summarized as:



(Change in one of the oracle components)



(smoother stabilizing controller)


(figure taken from Bouffard, et al)

# Experiments with LBMPC

- The second experiment involves ball catching.
- The (ping-pong sized) ball is thrown high in the air.
- The quadrotor need to catch it, slightly above ground (50cm).
- This is a challenging task, since the quadrotor has about 1 second to predict where the ball is going to land and to make the control decisions.
- The quadrotor continually updates the set point reference of the predicted landing point

# Ball-catching experiment

- The following illustrate one instance of this experiment:



(figure taken from Bouffard, et al)

- And the results are also available in vide:  
([https://www.youtube.com/watch?v=dL\\_ZFSvLXIU](https://www.youtube.com/watch?v=dL_ZFSvLXIU))

# Learning-Based Model Predictive Control (LBMPC)

- There is a technical detail about the LBMPC problem that is worth mentioning.
- Let's restate the problem again:

$$J_0(\bar{x}_0, \theta^{(0)}) = \min_{X, U, Z} \hat{J}_N(\tilde{x}_N) + \sum_{i=0}^{N-1} \tilde{x}_i^\top Q \tilde{x}_i + u_i^\top R u_i$$

$$\text{s.t. } \tilde{x}_{i+1} = A\tilde{x}_i + Bu_i + h(\tilde{x}_i, u_i; \theta^{(0)}), \quad \forall i \in \{0, 1, 2, \dots\}$$

$$\bar{x}_{i+1} = A\bar{x}_i + Bu_i, \quad \forall i \in \{0, 1, 2, \dots\}$$

$$u_i = K\bar{x}_i + z_i, \quad \forall i \in \{0, 1, 2, \dots\}$$

$$\bar{x}_i \in \mathcal{X} \ominus \mathcal{R}_i, \quad u_i \in \mathcal{U} \ominus (K \circ \mathcal{R}_k), \quad \forall i \in \{0, 1, 2, \dots\}$$

$$x_0 = \bar{x}_0$$

$$\bar{x}_N \in \mathcal{X}_f \ominus \mathcal{R}_N$$

# Learning-Based Model Predictive Control (LBMPC)

- Note that the optimization problem depends not only on the initial state  $\bar{x}_0$  but also on the oracle parameter vector  $\theta_0$ .
- In an abstract representation, we can write that problem as follows:

$$J_0(\bar{x}_0, \theta^{(0)}) = \min_{X, U, Z} F(X, U, Z)$$
$$\text{s.t. } (X, U, Z) \in G(\theta_0, \bar{x}_0)$$

- Where we highlight that the feasible region depends on  $\theta_0$ .



# Learning-Based Model Predictive Control (LBMPC)

- This is something we have encountered before, but never really addressed.
- Since our MPC algorithm is a **model-based** algorithm, one question to ask is whether we are able to learn in fact the true dynamics, if we use a “rich” enough oracle function.

- Namely suppose there exist a function  $h^*(x, u)$  such that:

$$x_{i+1} = Ax_i + Bu_i + h^*(x_i, u_i)$$

- And let  $J^*(\bar{x})$  be the Optimal value function if we solved the Optimal Control problem with  $h^*$  and starting from  $\bar{x}$ . Is it true that  $J(\bar{x}, \theta^{(t)}) \rightarrow J^*(\bar{x})$ , as  $t \rightarrow \infty$ ?

# Convergence of Approximate Optimization

- Recall the simple least squares problem where we wish to solve the following problem:

$$\theta^* = \arg \min_{\theta} \mathbb{E}[(y - x^\top \theta)^2]$$

- We typically cannot solve this problem due to the expectation. And we resort to solve, instead, a Sample Average Approximation (SAA):

$$\hat{\theta}_N = \arg \min_{\theta} \frac{1}{N} \sum_{i=1}^N (y_i - x_i^\top \theta)^2$$

- We are concerned if  $\hat{\theta}_N \rightarrow \theta^*$ .
- For this least-squares problem, this is true due to the Uniform Law of Large Numbers.

# Convergence of Approximate Optimization

- But in our LBMPC, we are solving a constrained optimization problem that changes over time, due to both initial condition and parameter vector.
- So it is not trivial that the Law of Large Numbers would hold in this case.
- Let  $F_n(x)$  be some function at stage  $n$ , with parameter vector  $\theta_n$ .
- We say that a function  $F_n$  epi-converges to another function  $F$  if and only if, at each point  $x$ :

$$\liminf_n F_n(x_n) \geq f(x), \forall x_n \rightarrow x$$

$$\limsup_n F_n(x_n) \leq f(x), \exists x_n \rightarrow x$$

# Convergence of Approximate Optimization

- This definition may be not intuitive. An intuitive explanation is say that  $F_n$  epi-converges to  $F$  if the epigraph of  $F_n$  converges to  $F$ .
- In addition if we minimize both functions over a bounded non-empty set  $X$ , it follows that:

$$V_n \rightarrow V \quad V_n = \min_{x \in X} \{F_n(x)\} \quad V = \min_{x \in X} \{F(x)\}$$

- And each approximated problem:  $V_n = \min_{x \in X} \{F_n(x)\}$
- Is feasible and form a bounded sequence where the set of optimal solutions will also converge. In the sense that:

$$\limsup_n (\arg \min F_n(x)) \subseteq \arg \min f(x)$$

# Convergence of LBMPC: Overview

- It turns out that we can apply this notion of epi-convergence to the LBMPC and prove that
- Using certain types of oracles such as:
  - Linear Oracles
  - Nadaraya-Watson (kernel-based) oracles
- We will converge, in the sense that the oracle will estimate accurately the model mismatches between reality and the nominal model.
- These proofs are a very technical and require a lot of groundwork.
- We present this very high-level view just to highlight the theoretical guarantees that LBMPC enjoy when employing function approximations.